



# Flight Fare Prediction

Welcome to our presentation on flight fare prediction. Flight ticket prices often change due to various factors like demand, season, and airline policies, making it hard for travelers to know the best time to book. This is where machine learning comes in. By analyzing past data such as routes, dates, airlines, and booking patterns, Data Science models can predict future prices with good accuracy. This helps travelers save money and make smarter booking decisions. It also benefits airlines by helping them optimize pricing and manage seat availability. In this presentation, we'll explore how this technology works and why it matters.

# *Domain and Context of Flight Fare Prediction*



## Airline Industry / Travel Tech

The prediction of flight ticket prices is a critical application within the airline industry and the broader travel technology sector. It supports strategic planning and operational efficiency for all participants.



## Dynamic Pricing & Planning

Flight fares are influenced by a multitude of factors including airline, origin, destination, journey date, duration, and number of stops. Accurate predictions empower customers to plan purchases effectively and allow companies to offer competitive, dynamic pricing strategies.



## Beneficiaries of Accurate Prediction

**Travelers:** Identify optimal booking times for cost savings.

**Travel Agencies:** Implement sophisticated dynamic pricing models.

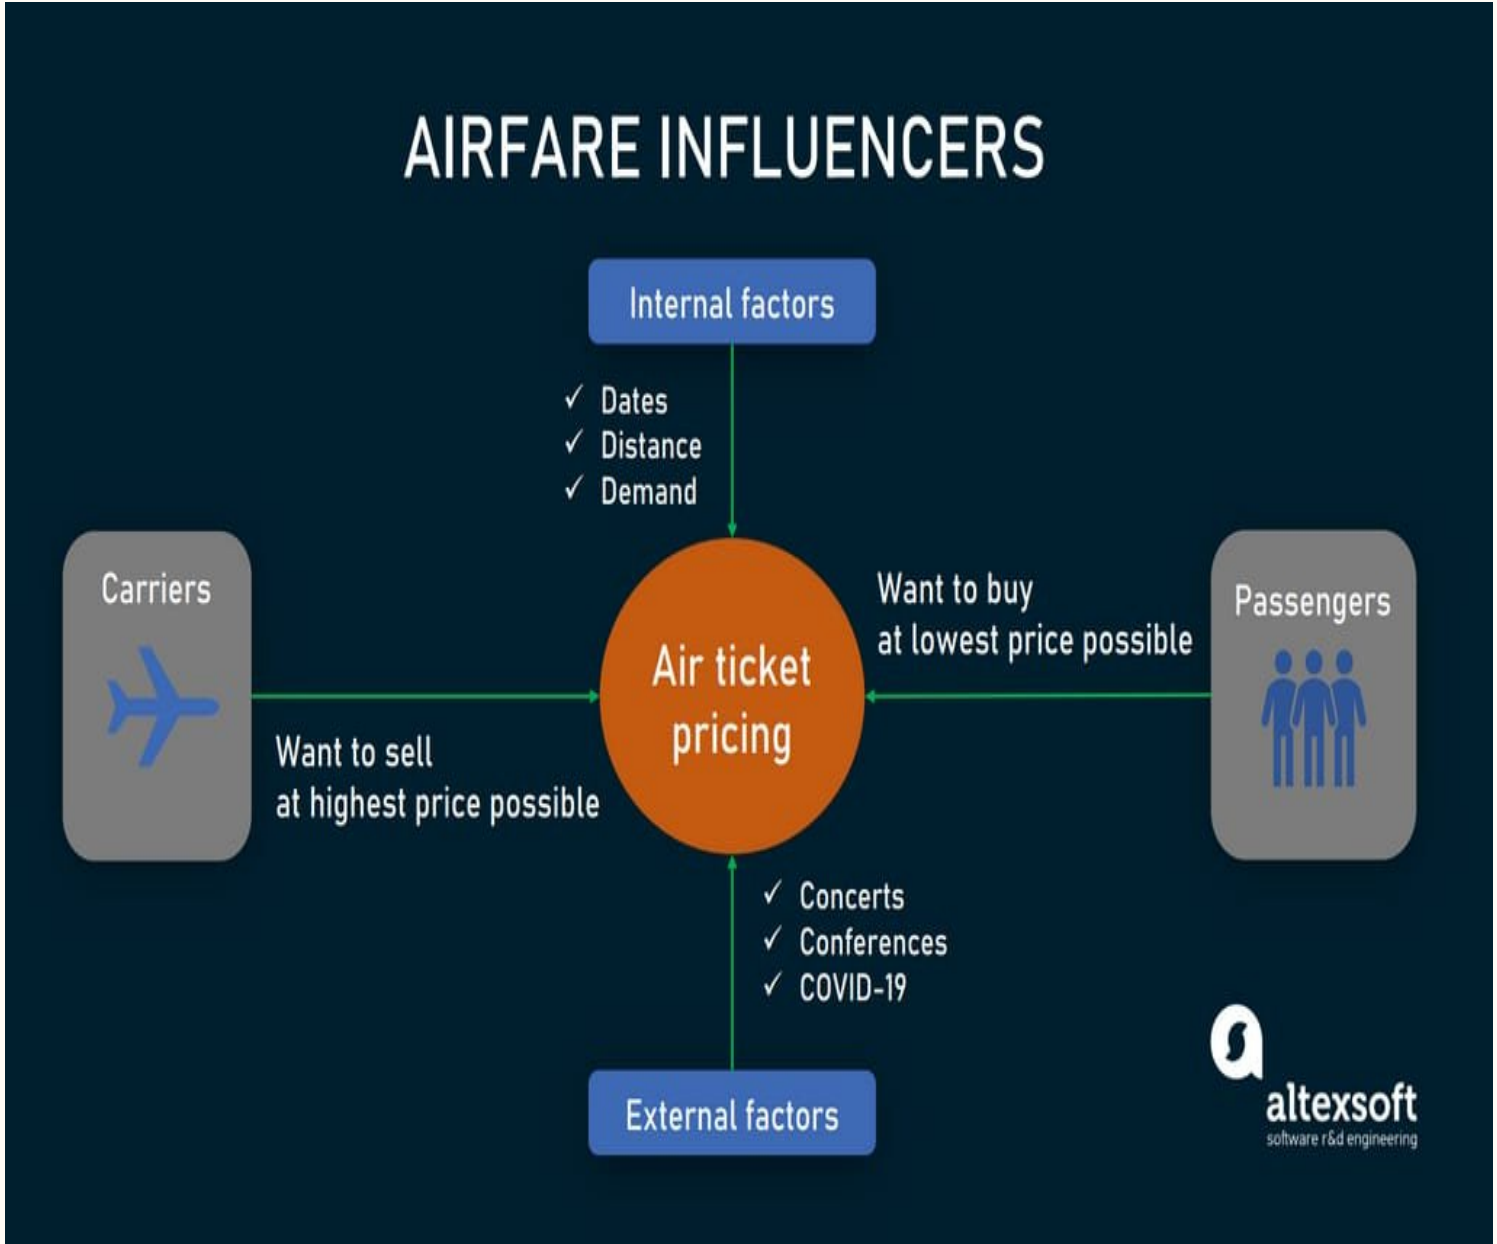
**Airlines:** Enhance demand forecasting and revenue management.

# Dataset Overview: Features and Categorization

The dataset used for this case study, often sourced from platforms like Kaggle, comprises various features crucial for fare prediction.

Airline	Categorical: Carrier of the flight.
Source, Destination	Categorical: Departure and arrival cities.
Date_of_Journey	Date/Time: Specific date of travel.
Duration	Textual/Mixed: Total travel time.
Total_Stops	Categorical: Number of layovers.
Additional_Info	Categorical: Miscellaneous flight details.
Price	Numerical: Target variable (flight fare).

Understanding these features and their types is fundamental for effective data preprocessing and model building.



# **Problem Statement & Workflow for Flight Fare Prediction**

## **Problem Statement:**

Airline ticket prices fluctuate based on various factors such as airline, source and destination cities, date of journey, duration, number of stops, and seasonality.

The goal of this project is to predict flight fares accurately so that:

- Travelers can make cost-effective booking decisions.
- Travel agencies/airlines can optimize their pricing strategies.

## **Work Done:**

1. Collected the flight fare dataset from Kaggle
2. Filled missing values using median/mode as required
3. Removed unnecessary columns from the dataset
4. Converted columns to correct data types
5. Scaled the numerical features for uniformity
6. Visualized relationships between key columns and flight fare

# Summary of Preprocessing Steps

## Handling Missing Values

Missing entries were imputed using appropriate statistical methods (e.g., **mean**, **median**, **mode**) to maintain data integrity.

## Duration Transformation

The 'Duration' feature, originally a string, was meticulously **converted into minutes** to allow for numerical analysis and model consumption.

## Date/Time Transformation

Granular features like **day**, **month**, and **hour** were extracted from 'Date\_of\_Journey', 'Dep\_Time', and 'Arrival\_Time' to capture temporal patterns.

## Categorical Data Encoding

Categorical features were transformed into numerical formats using **Label Encoding** or **One-Hot Encoding**, depending on their cardinality and potential ordinality.

## Outlier Removal

Outliers, particularly in the 'Price' distribution, were **identified and removed** to prevent undue influence on model training and improve prediction accuracy.

## Data Readiness

These comprehensive steps collectively ensured a **final cleaned dataset**, optimally prepared and ready for subsequent modelling phases.



## Raw Flight Dataset

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 13354 entries, 0 to 13353
Data columns (total 10 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Airline              13354 non-null  object
1   Date_of_Journey      13354 non-null  object
2   Source               13354 non-null  object
3   Destination          13354 non-null  object
4   Dep_Time             13354 non-null  object
5   Arrival_Time         13354 non-null  object
6   Duration             13354 non-null  object
7   Total_Stops          13353 non-null  object
8   Additional_Info      13354 non-null  object
9   is_train             13354 non-null  int64
dtypes: int64(1), object(9)
memory usage: 1.0+ MB
```

## Preprocessed Dataset with Encoded & Scaled Features

```
In [86]: X_train_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10683 entries, 0 to 10682
Data columns (total 33 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Total_Stops          10683 non-null  float64
1   Additional_Info      10683 non-null  float64
2   Date                 10683 non-null  float64
3   Month                10683 non-null  float64
4   Year                 10683 non-null  float64
5   Dep_hour             10683 non-null  float64
6   Dep_minute           10683 non-null  float64
7   Arrival_hour         10683 non-null  float64
8   Arrival_minute       10683 non-null  float64
9   Duration_minutes     10683 non-null  float64
10  Airline_Air Asia     10683 non-null  float64
11  Airline_Air India    10683 non-null  float64
12  Airline_GoAir        10683 non-null  float64
13  Airline_IndiGo       10683 non-null  float64
14  Airline_Jet Airways  10683 non-null  float64
15  Airline_Jet Airways Business  10683 non-null  float64
16  Airline_Multiple carriers  10683 non-null  float64
17  Airline_Multiple carriers Premium economy  10683 non-null  float64
18  Airline_SpiceJet     10683 non-null  float64
19  Airline_Trujet       10683 non-null  float64
20  Airline_Vistara       10683 non-null  float64
21  Airline_Vistara Premium economy  10683 non-null  float64
22  Source_Bangalore     10683 non-null  float64
23  Source_Chennai       10683 non-null  float64
24  Source_Delhi         10683 non-null  float64
25  Source_Kolkata       10683 non-null  float64
26  Source_Mumbai        10683 non-null  float64
27  Destination_Bangalore  10683 non-null  float64
28  Destination_Cochin   10683 non-null  float64
29  Destination_Delhi    10683 non-null  float64
30  Destination_Hyderabad  10683 non-null  float64
31  Destination_Kolkata   10683 non-null  float64
32  Destination_New Delhi  10683 non-null  float64
dtypes: float64(33)
memory usage: 2.7 MB
```

# Feature Engineering

```
#feature engineering
df['Date'] = df['Date_of_Journey'].str.split('/').str[0]
df['Month'] = df['Date_of_Journey'].str.split('/').str[1]
df['Year'] = df['Date_of_Journey'].str.split('/').str[2]
df.drop('Date_of_Journey',axis=1,inplace=True)
df.head()
```

```
df['Dep_hour'] = df['Dep_Time'].str.split(':').str[0]
df['Dep_minute'] = df['Dep_Time'].str.split(':').str[1]
df.drop('Dep_Time', axis=1, inplace=True)

df['Arrival_hour'] = df['Arrival_Time'].str.split(' ').str[0].str.split(':').str[0]
df['Arrival_minute'] = df['Arrival_Time'].str.split(' ').str[0].str.split(':').str[1]
df.drop('Arrival_Time', axis=1, inplace=True)

df.head()
```

```
#accessing unique values in total_stops column
df['Total_Stops'].unique()
```

```
array(['non-stop', '2 stops', '1 stop', '3 stops', nan, '4 stops'],
      dtype=object)
```

```
#converting text data to numerical values
df['Total_Stops'] = df['Total_Stops'].map({'non-stop':0, '2 stops':2, '1 stop':1, '3 stops':3, '4 stops':4})
df.head()
```

```
#splitting duration to hours and minutes
def extract_duration_hours(duration_str):
    try:
        if 'h' in duration_str:
            return int(duration_str.split('h')[0])
        else:
            return 0
    except:
        return 0

def extract_duration_minutes(duration_str):
    try:
        if 'm' in duration_str:
            parts = duration_str.split(' ')
            for part in parts:
                if 'm' in part:
                    return int(part.replace('m', ''))
            return 0
        else:
            return 0
    except:
        return 0
```

*# Apply the functions*

```
df['Duration_hour'] = df['Duration'].apply(extract_duration_hours)
df['Duration_minute'] = df['Duration'].apply(extract_duration_minutes)
df.drop('Duration', axis=1, inplace=True)
```

*#converting hours to minutes to reduce the features*

```
df['Duration_minutes'] = df['Duration_hour']*60 + df['Duration_minute']
df.drop(['Duration_hour','Duration_minute'], axis=1, inplace=True)
```

```
df.head()
```



# Original Dataset

First few rows of combined dataset:

```
Out[31]:
```

	Airline	Date_of_Journey	Source	Destination	Route	Dep_Time	Arrival_Time	Duration	Total_Stops	Additional_Info	is_train
0	IndiGo	24/03/2019	Banglore	New Delhi	BLR → DEL	22:20	01:10 22 Mar	2h 50m	non-stop	No info	1
1	Air India	1/05/2019	Kolkata	Banglore	CCU → IXR → BBI → BLR	05:50	13:15	7h 25m	2 stops	No info	1
2	Jet Airways	9/06/2019	Delhi	Cochin	DEL → LKO → BOM → COK	09:25	04:25 10 Jun	19h	2 stops	No info	1
3	IndiGo	12/05/2019	Kolkata	Banglore	CCU → NAG → BLR	18:05	23:30	5h 25m	1 stop	No info	1
4	IndiGo	01/03/2019	Banglore	New Delhi	BLR → NAG → DEL	16:50	21:35	4h 45m	1 stop	No info	1

```
#getting count of variuos data in 'additional info'  
df['Additional_Info'].value_counts()
```

```
Additional_Info  
No info          10493  
In-flight meal not included    2426  
No check-in baggage included    396  
1 Long layover          20  
Change airports          8  
Business class          5  
No Info          3  
1 Short layover          1  
Red-eye flight          1  
2 Long layover          1  
Name: count, dtype: int64
```

```
#dividing it to two categories  
df['Additional_Info'] = df['Additional_Info'].apply(lambda x: 'No info' if x == 'No info' else 'Others')  
df['Additional_Info'].value_counts()
```

```
Additional_Info  
No info    10493  
Others     2861  
Name: count, dtype: int64
```

# Encoding & Scaling

```
#encoding categorical columns  
%pip install scikit-learn  
from sklearn.preprocessing import LabelEncoder  
from sklearn.preprocessing import OneHotEncoder  
import pandas as pd  
  
le = LabelEncoder()  
df['Additional_Info'] = le.fit_transform(df['Additional_Info'])  
  
cols_to_encode = ['Airline', 'Source', 'Destination']  
  
ohe = OneHotEncoder(sparse_output=False, drop=None)  
encoded = ohe.fit_transform(df[cols_to_encode])  
encoded_cols = ohe.get_feature_names_out(cols_to_encode)  
encoded_df = pd.DataFrame(encoded, columns=encoded_cols, index=df.index)  
df = pd.concat([df.drop(cols_to_encode, axis=1), encoded_df], axis=1)
```

```
#scale numeric features to mean 0 and std dev 1  
from sklearn.preprocessing import StandardScaler  
import pandas as pd
```

```
#splitting dataset back to training and testing using flag  
train_clean = df[df['is_train'] == 1].drop('is_train', axis=1)  
test_clean = df[df['is_train'] == 0].drop('is_train', axis=1)
```

```
scaler = StandardScaler()  
X_train = scaler.fit_transform(train_clean)  
X_test = scaler.transform(test_clean)  
  
X_train_df = pd.DataFrame(X_train, columns=train_clean.columns)  
X_test_df = pd.DataFrame(X_test, columns=test_clean.columns)
```

```
X_train_df.head()
```

	Total_Stops	Additional_Info	Date	Month	Year	Dep_hour	Dep_minute	Arrival_hour	Arrival_minute	Duration_minutes	...
0	-1.220744	-0.529309	1.237383	-1.467490	0.0	1.654259	-0.235050	-1.800427	-0.890057	-0.931583	...
1	1.741483	-0.529309	-1.475239	0.250276	0.0	-1.303095	1.363492	-0.050851	-0.587124	-0.390072	...
2	1.741483	-0.529309	-0.531719	1.109160	0.0	-0.607247	0.031373	-1.363033	0.018744	0.978475	...
3	0.260370	-0.529309	-0.177898	0.250276	0.0	0.958411	-1.034321	1.407129	0.321677	-0.626367	...
4	0.260370	-0.529309	-1.475239	-1.467490	0.0	0.610487	1.363492	1.115533	0.624611	-0.705132	...


5 rows × 33 columns



# Visualization of Applied Transformation

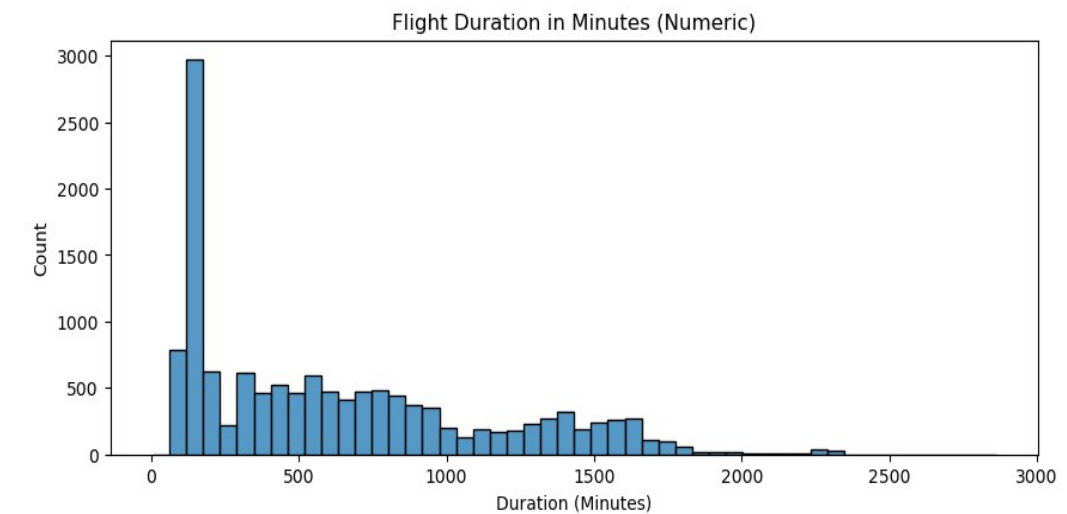
Plot for Total Stops Distribution after Encoding

```
In [59]: plt.figure(figsize=(10,4))
sns.countplot(x='Total_Stops', data=df)
plt.title("Total Stops Distribution After Transformation")
plt.xlabel("Stops")
plt.ylabel("Count")
plt.show()
```



Plot for Flight Distribution After Conversion to Minutes

```
[127]: plt.figure(figsize=(10,4))
sns.histplot(df['Duration_minutes'], bins=50)
plt.title("Flight Duration in Minutes (Numeric)")
plt.xlabel("Duration (Minutes)")
plt.ylabel("Count")
plt.show()
```



# **Model Implementation: K-Nearest Neighbours (KNN) Classifier**

- **Objective:** To classify flight fares into three categories: 'Low', 'Medium', and 'High'.

## ***# Convert the target variable 'Price' into classes***

```
price_bins = [0, 6000, 12000, float('inf')]
```

```
price_labels = ['Low', 'Medium', 'High']
```

```
y_class = pd.cut(y, bins=price_bins, labels=price_labels, right=False)
```

## ***# Import the classifier and split the data for validation***

```
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn.model_selection import train_test_split
```

## ***# Split the preprocessed training data for training and validation***

```
X_train, X_val, y_train_class, y_val_class = train_test_split(
```

```
X_train_df, y_class, test_size=0.2, random_state=42, stratify=y_class)
```

## ***# Initialize and train the KNN Classifier with K=5***

```
knn_classifier = KNeighborsClassifier(n_neighbors=5)
```

```
knn_classifier.fit(X_train, y_train_class)
```

### **Target Variable Distribution:**

- Medium: 4620
- Low: 3167
- High: 2896

- 1. The trained KNN model was evaluated on the validation set to assess its performance.

**Model Evaluation**

- **2. Accuracy Score:** The model achieved an accuracy of **86%** on the validation data.

- **3. Confusion Matrix:**

[ [485 7 87]  
[ 2 564 68]  
[ 71 66 787] ]

Classification Report Structure				
	precision	recall	f1-score	support
Class 1 (e.g., High)	0.87	0.84	0.85	579
Class 2 (e.g., Low)	0.89	0.89	0.89	634
Class 3 (e.g., Medium)	0.84	0.85	0.84	924
accuracy			0.86	2137
macro avg	0.86	0.86	0.86	2137
weighted avg	0.86	0.86	0.86	2137



# ***Conclusion & Insights***

- ***Summary:***
- This project successfully developed a machine learning model to predict flight fare categories with a high degree of accuracy.
- Extensive data preprocessing and feature engineering were performed to prepare the dataset for modeling. This included handling missing values, converting data types, and encoding categorical variables.
- A K-Nearest Neighbors (KNN) classification model was implemented and trained on the cleaned dataset.
- ***Key Insights:***
- The model achieved an accuracy of 86%, demonstrating its effectiveness in classifying flight prices.
- The classification report shows strong precision and recall across all price categories, indicating a well-balanced model.
- Hyperparameter tuning helped in identifying the optimal K value, which is crucial for the performance of the KNN algorithm.